

Verifying filesystems in ACL2

Project report, CS380L Fall 2017

Mihir Mehta

Department of Computer Science
University of Texas at Austin

`mihir@cs.utexas.edu`

07 December, 2017

Outline

Motivation

Our approach

Progress so far

Outline

Motivation

Our approach

Progress so far

Why we need a verified filesystem

- ▶ Modern filesystems have become increasingly complex, and so have the tools to analyse and recover data from them.
- ▶ It would be worthwhile to specify and formally verify, in the ACL2 theorem prover, the guarantees claimed by filesystems and tools.
- ▶ Work on this project started last year - since then I've built 5 increasingly complex models.
- ▶ This talk will be focussed on the 4th and 5th models.


Outline

Motivation

Our approach

Progress so far

File system models




Minimal set of operations?

- ▶ The Google filesystem suggests a minimal set of operations:
 - ▶ create
 - ▶ delete
 - ▶ open
 - ▶ close
 - ▶ read
 - ▶ write
- ▶ Of these, open and close require the maintenance of file descriptor state - so they can wait.
- ▶ However, they are essential when describing concurrency and multiprogramming behaviour.
- ▶ Thus, we can start modelling a filesystem, and several refinements thereof.

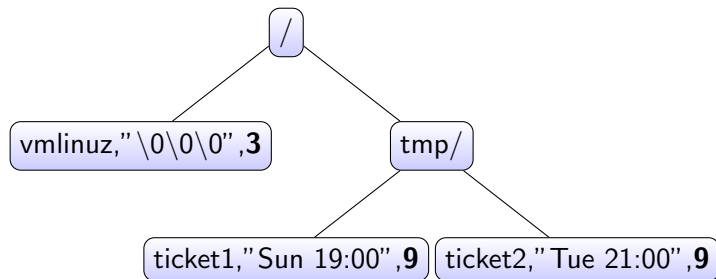
Quick overview of models

- ▶ What follows is a sequence of **increasingly refined** models.
- ▶ Model 1: Tree representation of directory structure with unbounded file size and unbounded filesystem size.
- ▶ Model 2: Model 1 with file length as metadata.
- ▶ Model 3: Tree representation of directory structure with file contents stored in a "disk" (unbounded in length).
- ▶ Model 4: Model 3 with bounded filesystem size and garbage collection.
- ▶ Model 5: Model 4 with permissions for read/write for the user and others (no groups as yet)

Model 1



Model 2



Model 3




Table: Disk

0	\0\0\0
1	Sun 19:0
2	0
3	Tue 21:0
4	0

Model 4




Table: Disk and allocation vector

0	\0\0\0	true
1	Sun 19:0	true
2	0	true
3	Tue 21:0	true
4	0	true
5		false

Outline

Motivation

Our approach

Progress so far

Proof approaches and techniques

- ▶ There are many properties that could be considered for correctness, but we choose to focus on the read-over-write theorems from the first-order theory of arrays.
- ▶ Read n characters starting at position $start$ in the file at path hns in filesystem fs :
`l1-rdchs(hns, fs, start, n)`
- ▶ Write string $text$ characters starting at position $start$ in the file at path hns in filesystem fs :
`l1-wrchs(hns, fs, start, text)`

Proof approaches and techniques

- ▶ The first read-over-write theorem defines the semantics of reading from a location after writing to the same location. Formally, assuming $n = \text{length}(\text{text})$ and suitable "type" hypotheses (omitted here):

$$\begin{aligned} & \text{l1-rdchs}(\text{hns}, \text{l1-wrchs}(\text{hns}, \text{fs}, \text{start}, \text{text}), \\ & \quad \text{start}, n) = \\ & \text{text} \end{aligned}$$

- ▶ The second read-over-write theorem defines the semantics of reading from a location after writing to a different location. Formally, assuming $\text{hns1} \neq \text{hns2}$ and suitable "type" hypotheses (omitted here):

$$\begin{aligned} & \text{l1-rdchs}(\text{hns1}, \text{l1-wrchs}(\text{hns2}, \text{fs}, \text{start2}, \text{text2}), \\ & \quad \text{start1}, n1) = \\ & \text{l1-rdchs}(\text{hns1}, \text{fs}, \text{start1}, n1) \end{aligned}$$

Proof approaches and techniques

- ▶ For each of the models 1 through 5, we have proofs of correctness of the two read-after-write properties, making use of the proofs of equivalence between models and their successors.
- ▶ For model 5, the proof assumes the permissions are granted.
- ▶ The proof of the converse property - that reads and writes fail when permissions are not granted - remains to be done.

Proof example: first read-over-write in model 2

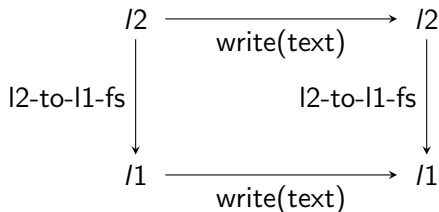


Figure: I2-wrchs-correctness-1

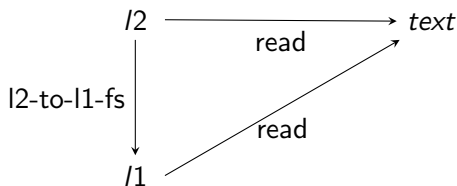


Figure: I2-rdchs-correctness-1

Proof example: first read-over-write in model 2




Figure: l2-read-over-write-1

Source analysis

Table: Code written for this project

Source lines of code (ACL2)	4964
defun events (function definitions)	106
defthm events (lemmas and proofs)	374