

# Verifying filesystems in ACL2

Project report, CS380L Fall 2017

Mihir Mehta

Department of Computer Science  
University of Texas at Austin


`mihir@cs.utexas.edu`

07 December, 2017

# Why we need a verified filesystem

- ▶ Modern filesystems have become increasingly complex, and so have the tools to analyse and recover data from them.
- ▶ It would be worthwhile to specify and formally verify, in the ACL2 theorem prover, the guarantees claimed by filesystems and tools.
- ▶ Work on this project started last year - since then I've built 5 increasingly complex models.
- ▶ The plan is to model FAT32, adding features in every model and maintaining proofs of correctness.

# File system models



# Minimal set of operations?

- ▶ The Google filesystem suggests a minimal set of operations:
  - ▶ create
  - ▶ delete
  - ▶ open
  - ▶ close
  - ▶ read
  - ▶ write
- ▶ Of these, open and close require the maintenance of file descriptor state - so they can wait.
- ▶ However, they are essential when describing concurrency and multiprogramming behaviour.
- ▶ Thus, we can start modelling a filesystem, and several refinements thereof.

## More about the models

- ▶ Model 1: Tree representation of directory structure with unbounded file size and unbounded filesystem size.
- ▶ Model 2: Model 1 with file length as metadata.
- ▶ Model 3: Tree representation of directory structure with file contents stored in a "disk" (unbounded in length).
- ▶ Model 4: Model 3 with bounded filesystem size and garbage collection.
- ▶ Model 5: Model 4 with permissions for read/write for the user and others (no groups as yet)

# Proof approaches and techniques

- ▶ There are many properties that could be considered for correctness, but we choose to focus on the read-over-write theorems from the first-order theory of arrays.
- ▶ Read  $n$  characters starting at position  $start$  in the file at path  $hns$  in filesystem  $fs$ :  
`l1-rdchs(hns, fs, start, n)`
- ▶ Write string  $text$  characters starting at position  $start$  in the file at path  $hns$  in filesystem  $fs$ :  
`l1-wrchs(hns, fs, start, text)`

# Proof approaches and techniques

- ▶ The first read-over-write theorem defines the semantics of reading from a location after writing to the same location. Formally, assuming  $n = \text{length}(\text{text})$  and suitable "type" hypotheses (omitted here):

$$\begin{aligned} & \text{ll-rdchs}(\text{hns}, \text{ll-wrchs}(\text{hns}, \text{fs}, \text{start}, \text{text}), \\ & \quad \text{start}, n) = \\ & \text{text} \end{aligned}$$

- ▶ The second read-over-write theorem defines the semantics of reading from a location after writing to a different location. Formally, assuming  $\text{hns1} \neq \text{hns2}$  and suitable "type" hypotheses (omitted here):

$$\begin{aligned} & \text{ll-rdchs}(\text{hns1}, \text{ll-wrchs}(\text{hns2}, \text{fs}, \text{start2}, \text{text2}), \\ & \quad \text{start1}, n1) = \\ & \text{ll-rdchs}(\text{hns1}, \text{fs}, \text{start1}, n1) \end{aligned}$$

# Proof approaches and techniques

- ▶ For each of the models 1 through 5, we have proofs of correctness of the two read-after-write properties, making use of the proofs of equivalence between models and their successors.
- ▶ For model 5, the proof assumes the permissions are granted.
- ▶ The proof of the converse property - that reads and writes fail when permissions are not granted - remains to be done.



## Proof example: first read-over-write in model 2

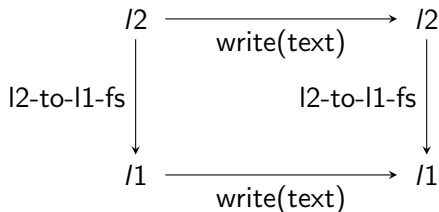


Figure: l2-wrchs-correctness-1

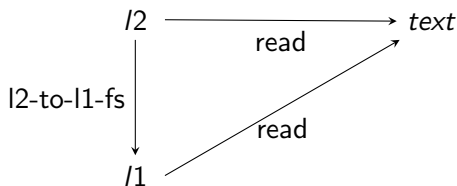


Figure: l2-rdchs-correctness-1

## Proof example: first read-over-write in model 2

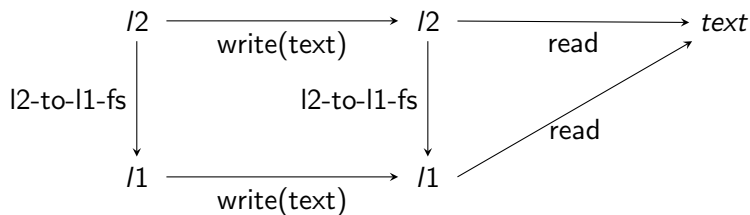


Figure: l2-read-over-write-1

# Source analysis

Table: Code written for this project

Source lines of code (ACL2)	6017
defun events (function definitions)	118
defthm events (lemmas and proofs)	419