

# Verifying the FAT32 filesystem in ACL2

Mihir Mehta

Department of Computer Science  
University of Texas at Austin

`mihir@cs.utexas.edu`

06 April, 2018

# Outline

Motivation and related work

Our approach

# Outline

Motivation and related work

Our approach

# Why we need a verified filesystem

- ▶ Filesystems are everywhere, even as operating systems move towards making them invisible.
- ▶ In the absence of a clear specification of filesystems, users (and sysadmins in particular) are underserved.
- ▶ Modern filesystems have become increasingly complex, and so have the tools to analyse and recover data from them.
- ▶ It would be worthwhile to specify and formally verify, in the ACL2 theorem prover, the guarantees claimed by filesystems and tools.

## Related work

- ▶ In Haogang Chen's 2016 dissertation, the author uses Coq to build a filesystem (named FSCQ) which is proven safe against crashes in a new logical framework named Crash Hoare Logic. His (exported) Haskell implementation performs comparably to ext4.
- ▶ Hyperkernel (Nelson et al., SOSP '17) is a "push-button" verification effort, but approximates by changing POSIX system calls for ease of verification.
- ▶ In our work, we instead aim to model an existing filesystem (FAT32) faithfully and match the resulting disk image byte-to-byte.

# Outline

Motivation and related work

Our approach

# Modelling a filesystem

- ▶ We opt to iteratively model a filesystem, incrementally adding features of FAT32.
- ▶ This also allows us to prove correctness in an iterative fashion, by proving equivalences and thereby reusing correctness results from previous models.
- ▶ Apart from noting that some of these models were based on the previous technology target, the CP/M filesystem, we will not dwell on these.

# Modelling a filesystem

- ▶ In our most recent model, we separate our filesystem into:
  - ▶ a tree, in which non-leaf nodes represent (sub)directories and leaf nodes represent regular files;
  - ▶ a disk, containing the textual contents of regular files broken into fixed-size blocks;
  - ▶ and a file allocation table, mapping each block in a regular file to the next, this allowing us to read the contents of the entire file.



# Verifying the model

- ▶ We've focussed so far on two filesystem properties, derived from the theory of arrays where they are known as the read-over-write properties.
  1. After a write, a read of the same length at the same location should yield that which was written.
  2. After a write, a read at a different location should yield the same results as a read before the write.
- ▶ These properties have been proven for all models so far, including the present model which features a file allocation table.

# Proof challenges

- ▶ How do we define a "good state" of a filesystem, which shows that reading, writing and other operations can be safely carried out?
- ▶ Answering this question involves a trade-off between simplicity (to help with verification) and generality (to model as many real-world situations as possible.)
- ▶ We choose to require:
  - ▶ that each block on the disk is attributed to at most one regular file;
  - ▶ that the clusters attributed to each non-empty regular file end with a legal EOF value, as defined by the FAT specification.
  - ▶ that each regular file is annotated with "length", a metadata field that corresponds to the actual length of the file as determined by traversing the file allocation table and reading the corresponding blocks.

# Validating the model

# Future work

1. Complete the FAT32 model, by means of
  - ▶ supporting variable cluster sizes,
  - ▶ moving the file allocation table onto the disk, and
  - ▶ moving all file and directory metadata from the tree to the disk.
2. Model a more complex filesystem, for instance NTFS, by re-using algorithms and proofs from the models built so far.