

Verifying filesystems in ACL2

Towards verifying file recovery tools

Mihir Mehta

Department of Computer Science
University of Texas at Austin

`mihir@cs.utexas.edu`

14 November, 2017

Outline

Motivation and related work

Our approach

Progress so far

Future work

Why we need a verified filesystem

- ▶ Filesystems are everywhere, even as operating systems move towards making them invisible.
- ▶ In the absence of a clear specification of filesystems, users (and sysadmins in particular) are underserved.
- ▶ Modern filesystems have become increasingly complex, and so have the tools to analyse and recover data from them.
- ▶ It would be worthwhile to specify and formally verify, in the ACL2 theorem prover, the guarantees claimed by filesystems and tools.

Related work

- ▶ In Haogang Chen's 2016 dissertation, the author uses Coq to build a filesystem (named FSCQ) which is proven safe against crashes in a new logical framework named Crash Hoare Logic. His (exported) Haskell implementation performs comparably to ext4.
- ▶ Hyperkernel (Nelson et al., SOSP '17) is a "push-button" verification effort, but approximates by changing POSIX system calls for ease of verification.
- ▶ In our work, we instead aim to model an existing filesystem (FAT32) faithfully and match the resulting disk image byte-to-byte.

Outline

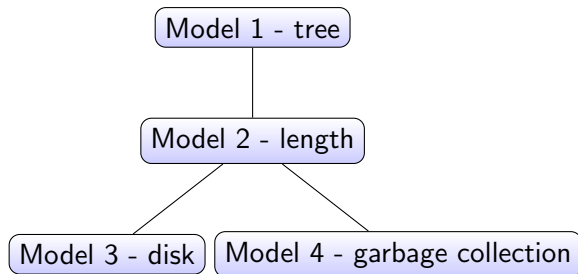
Motivation and related work

Our approach

Progress so far

Future work

File system models



Choosing an initial model

- ▶ Our goal here is to verify the FAT32 filesystem, but we need a simpler model to begin with.
- ▶ Our filesystem's operations should suffice for running a workload.
- ▶ Yet, parsimony and avoidance of redundancy are helpful for theorem proving.
- ▶ What's a necessary and sufficient set of operations?

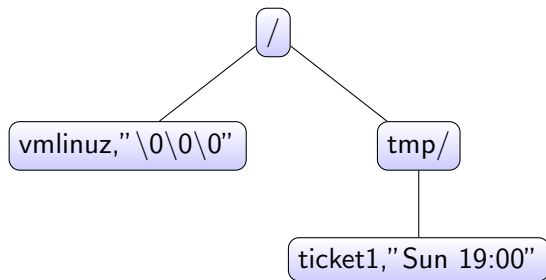
Minimal set of operations?

- ▶ The Google filesystem suggests a minimal set of operations:
 - ▶ create
 - ▶ delete
 - ▶ open
 - ▶ close
 - ▶ read
 - ▶ write
- ▶ Of these, open and close require the maintenance of file descriptor state - so they can wait.
- ▶ However, they are essential when describing concurrency and multiprogramming behaviour.
- ▶ Thus, we can start modelling a filesystem, and several refinements thereof.

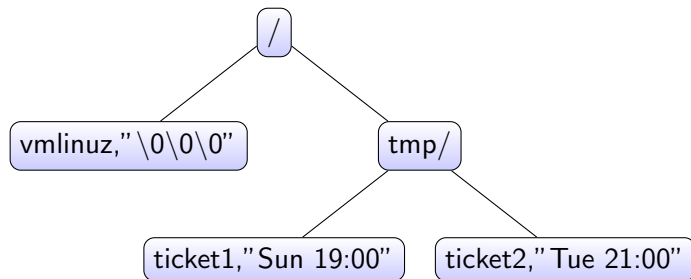
Quick overview of models

- ▶ What follows is a sequence of **increasingly refined** models.
- ▶ Model 1: Tree representation of directory structure with unbounded file size and unbounded filesystem size.
- ▶ Model 2: Model 1 with file length as metadata.
- ▶ Model 3: Tree representation of directory structure with file contents stored in a "disk" (unbounded in length).
- ▶ Model 4: Model 3 with bounded filesystem size and garbage collection.
- ▶ In our running example, we'll observe file creation, deletion, and overwriting in each of these models in turn.

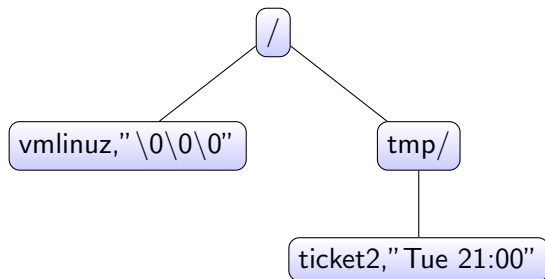
Model 1




Model 1



Model 1



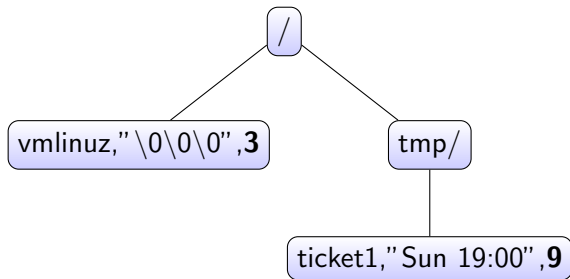
Model 1



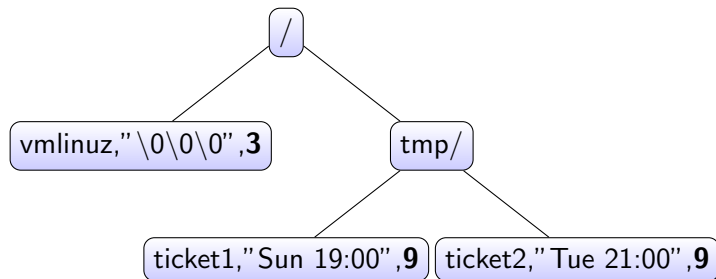
Model 2

- ▶ Model 1 supports nested directory structures, unbounded file size and unbounded filesystem size.
- ▶ However, there's no metadata, either to provide additional information or to validate the contents of the file.
- ▶ With an extra field for length, we can create a simple version of fsck that checks file contents for consistency.
- ▶ Further, we can verify that create, write, delete etc preserve this notion of consistency.

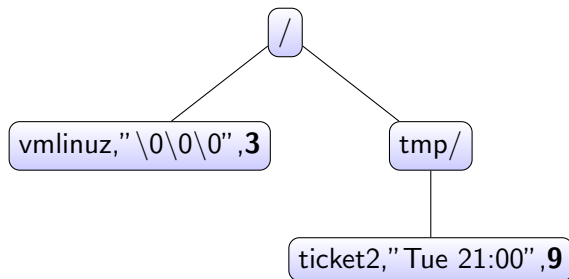
Model 2




Model 2



Model 2



Model 2



Model 3

- ▶ As the next step, we focus on externalising the storage of file contents.
- ▶ We also choose to break up file contents into "blocks" of a constant length (8).
 - ▶ Note: this would mean storing file length is no longer optional, to avoid reading garbage past end of file at the end of a block.

Model 3

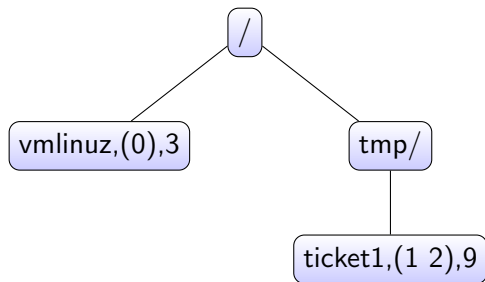


Table: Disk

0	\0\0\0
1	Sun 19:0
2	0

Model 3

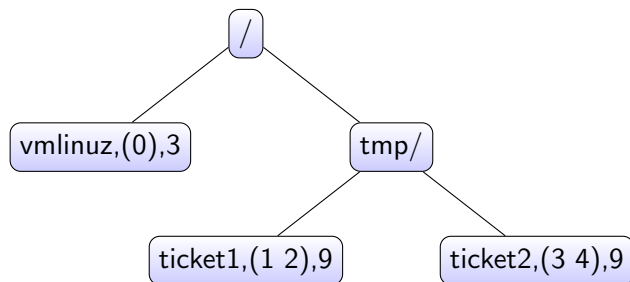


Table: Disk

0	\0\0\0
1	Sun 19:0
2	0
3	Tue 21:0
4	0

Model 3

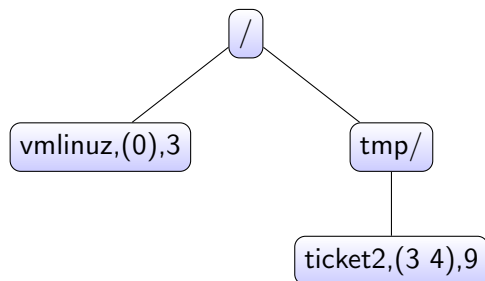


Table: Disk

0	\0\0\0
1	Sun 19:0
2	0
3	Tue 21:0
4	0

Model 3

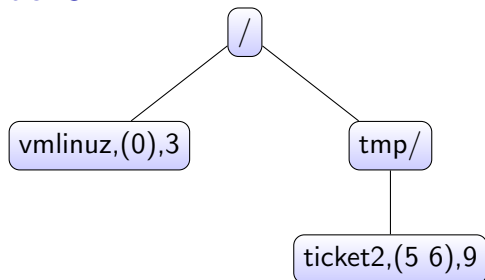


Table: Disk

0	\0\0\0
1	Sun 19:0
2	0
3	Tue 21:0
4	0
5	Wed 01:0
6	0

Model 4

- ▶ In the fourth model, we attempt to implement garbage collection in the form of an allocation vector.
- ▶ The allocation vector tracks whether blocks in the filesystem are in use by a file. This allows us to reuse unused blocks.

Model 4




Table: Disk and allocation vector

0	\\0\\0\\0	true
1	Sun 19:0	true
2	0	true
3		false
4		false
5		false

Model 4

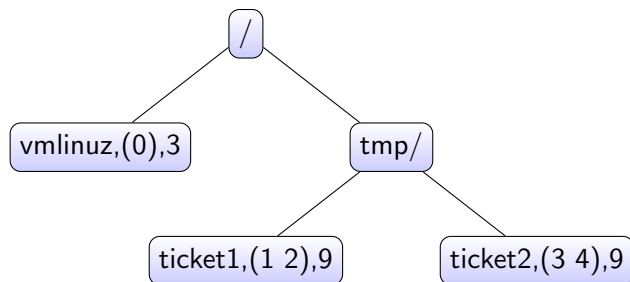


Table: Disk and allocation vector

0	\0\0\0	true
1	Sun 19:0	true
2	0	true
3	Tue 21:0	true
4	0	true
5		false

Model 4




Table: Disk and allocation vector

0	\\0\\0\\0	true
1	Sun 19:0	false
2	0	false
3	Tue 21:0	true
4	0	true
5		false

Model 4

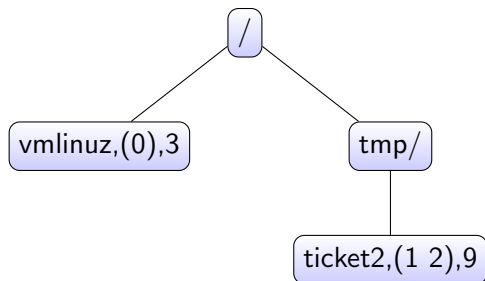


Table: Disk and allocation vector

0	\0\0\0	true
1	Wed 01:0	true
2	0	true
3	Tue 21:0	false
4	0	false
5		false

Outline

Motivation and related work

Our approach

Progress so far

Future work

Proof approaches and techniques

- ▶ There are many properties that could be considered for correctness, but we choose to focus on the read-over-write theorems from the first-order theory of arrays.
- ▶ Read n characters starting at position $start$ in the file at path hns in filesystem fs :
`l1-rdchs(hns, fs, start, n)`
- ▶ Write string $text$ characters starting at position $start$ in the file at path hns in filesystem fs :
`l1-wrchs(hns, fs, start, text)`

Proof approaches and techniques

- ▶ First read-over-write theorem: reading from a location after writing to the same location should yield the data that was written. Formally, assuming $n = \text{length}(\text{text})$ and suitable "type" hypotheses (omitted here):

```
l1-rdchs(hns, l1-wrchs(hns, fs, start, text),  
start, n)  
=  
text
```

- ▶ Second read-over-write-theorem: Reading from a location after writing to a different location should yield the same result as reading before writing. Formally, assuming $\text{hns1} \neq \text{hns2}$ and suitable "type" hypotheses (omitted here):

```
l1-rdchs(hns1, l1-wrchs(hns2, fs, start2, text2),  
start1, n1)  
=  
l1-rdchs(hns1, fs, start1, n1)
```

Proof approaches and techniques

- ▶ For each of the models 1, 2, 3 and 4, we have proofs of correctness of the two read-after-write properties, making use of the proofs of equivalence between models and their successors.
- ▶ Model 4 presented some unique challenges - proving the read-after-write properties required proving an equivalence between model 4 and model 2, rather than model 3.

Proof example: read-over-write in model 2

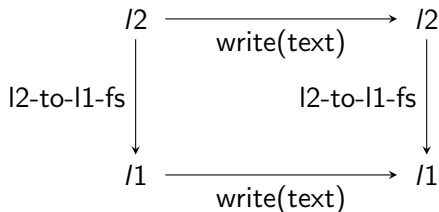


Figure: l2-wrchs-correctness-1

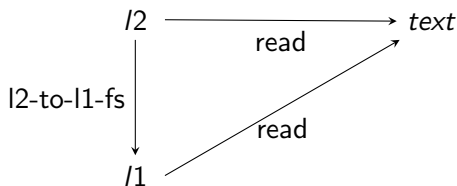


Figure: l2-rdchs-correctness-1

Proof example: read-over-write in model 2

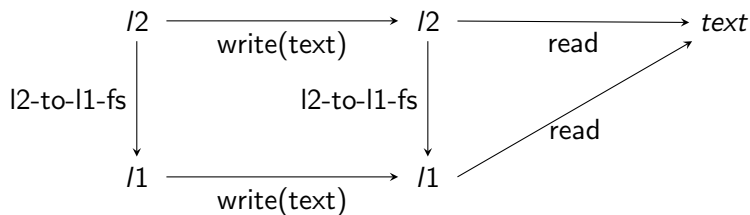


Figure: l2-read-over-write-1

Source analysis

Table: Code (over 3 semesters)

sloccount (lines of ACL2 code)	4964
defun events (function definitions)	106
defthm events (lemmas and proofs)	374

Outline

Motivation and related work

Our approach

Progress so far

Future work

Future work

- ▶ Model and verify file permissions.
- ▶ Linearise the tree, leaving only the disk.
- ▶ Add the system call open and close with the introduction of file descriptors.
This would be a step towards the study of concurrent FS operations.
- ▶ Eventually emulate the FAT32 filesystem as a convincing proof of concept, and move on to fsck and file recovery tools.